

# Stochastic Models: Model Selection and Optimisation

Lecture 6 – From Evaluation to Decisions

Sebastian Müller

Lecture 6



# Where We Are

---

- ▶ **Lecture 1:** Motivation, Poisson refresher, simulation mindset.
- ▶ **Lecture 2:** Queueing basics, M/M/1 and M/M/c, Little's Law.
- ▶ **Lecture 3:** Error control, CLT and concentration, variance reduction, sensitivity.
- ▶ **Lecture 4:** Parameter estimation for Poisson, M/M/1, M/G/1; plug-in performance.
- ▶ **Lecture 5:** M/G/2 evaluation exercise from event log to fitted model.

*Today: move from evaluation of a given system to optimisation of design choices.*

# Goals for Lecture 6

---

- ▶ Decide which queueing model ( $G/G/c/K/\dots$ ) is appropriate for a given dataset and question.
- ▶ Use fitted models to run simulation experiments with confidence intervals.
- ▶ Perform **scenario analysis**: how changes in load, variability, or capacity affect performance.
- ▶ Introduce **costs** and formulate a simple optimisation problem over design variables (e.g. number of servers).
- ▶ Connect the full pipeline: data → model → estimation → simulation → decision.

# From Data to Candidate Models

---

- ▶ Starting point: cleaned event log (Lecture 5) with arrivals, service times, and queue lengths.
- ▶ We want a **parsimonious** model that captures key features relevant for decisions.
- ▶ Two equally important dimensions:
  - ▶ **Structural choice:** M/M/1, M/M/c, M/G/1, M/G/c, GI/G/c, finite buffers, abandonment, priorities, networks.
  - ▶ **Parameter estimation:**  $\hat{\lambda}$ , service distribution parameters, variability measures.
- ▶ Model choice should be driven by *data, domain knowledge, and decision needs*.

# Dimensions of Choice

---

- ▶ **Arrival process**

- ▶ Poisson (homogeneous) vs. time-varying rate vs. overdispersed arrivals.
- ▶ Diagnostics: inter-arrival histogram, ECDF, overdispersion of counts, time-of-day effects.

- ▶ **Service-time distribution**

- ▶ Exponential vs. deterministic offset + exponential vs. heavy-tailed.
- ▶ Diagnostics: empirical CV, skewness, log-survival plot, QQ-plot against exponential.

- ▶ **System structure**

- ▶ Number of servers  $c$ , buffer capacity  $K$ , priority classes, reneging/abandonment.
- ▶ Single-node approximation vs. network / multi-stage system.

# Simple vs Complex Models

---

## Idea

Start with a simple baseline model and add complexity only when necessary for decisions.

- ▶ Baselines:
  - ▶ M/M/c for systems with no strong evidence against exponential assumptions.
  - ▶ M/G/1 or M/G/c when service-time variability clearly deviates from exponential.
- ▶ More flexible options:
  - ▶ GI/G/c approximations when inter-arrivals also deviate from Poisson.
  - ▶ Empirical service distributions (resampling) when no simple parametric family fits well.
- ▶ Trade-off:
  - ▶ Simpler models are easier to explain and calibrate.
  - ▶ Richer models may better capture tails or bursts.

# Model Selection Guidelines

---

- ▶ **Structural consistency:** data must not contradict assumptions (e.g. observed concurrency  $\leq c$  if you model  $c$  servers; no abandonment if model excludes it).
- ▶ **Descriptive fit:** do simulated histograms / quantiles for arrivals and services match the data?
- ▶ **Performance fit:** does the model reproduce key metrics (mean  $W_q$ ,  $L_q$ , utilisation) within uncertainty bands?
- ▶ **Decision relevance:** only add complexity if it changes the recommendation (e.g. number of servers to deploy).
- ▶ **Interpretability:** stakeholders should understand the assumptions and their limits.

*Pick the simplest model that fits the data and supports the decision without violating observed behaviour.*

# Case Study Pitfall: Fitting M/G/3 Data with M/G/2

---

- ▶ **What went wrong in Lecture 5?**

- ▶ Data inspection showed moments with 3 jobs in service  $\Rightarrow$  evidence for  $c = 3$ .
- ▶ Fitting M/G/2 underestimates capacity, inflates utilisation, and biases wait predictions.

- ▶ **How to avoid it**

- ▶ Plot max number in service over time; compare to assumed  $c$ .
- ▶ Refit with  $c = 2$  vs  $c = 3$  and compare simulated  $W_q$ ,  $L_q$ , tails *with CIs*.
- ▶ Prefer the model that (i) matches observed service concurrency and (ii) does not degrade fit on waits/queues.

- ▶ **Report explicitly** when structural choices are ambiguous; give a decision recommendation for each plausible  $c$ .

# What-If Analysis

---

## Idea

Use the fitted model as a baseline and explore how performance changes when we perturb inputs or design choices.

- ▶ Vary arrival rate:  $\lambda = \alpha \hat{\lambda}$  for  $\alpha \in \{0.8, 1.0, 1.2\}$  (demand uncertainty).
- ▶ Vary number of servers:  $c \in \{1, 2, 3, 4\}$  (staffing / capacity decisions).
- ▶ Optionally vary service variability: more/less variable  $G$  with the same mean.
- ▶ For each scenario, simulate and compute CIs for  $W_q$ ,  $L_q$ , utilisation, and service-level metrics.

*Goal: identify regimes where the system is robust and regimes where it is fragile (near saturation).*

# Interpreting Scenario Tables

---

- ▶ Present results as tables or heatmaps:
  - ▶ Rows: number of servers  $c$ .
  - ▶ Columns: demand multiplier  $\alpha$  or other scenario parameter.
- ▶ For each cell, report:
  - ▶ Point estimates and 95% CIs for  $W_q$ ,  $L_q$ , utilisation.
  - ▶ Possibly probability of violating a service-level target (e.g.  $P(W_q > w^*)$ ).
- ▶ Use these tables to communicate trade-offs:
  - ▶ Where does adding a server significantly improve waiting times?
  - ▶ Where do CIs for different configurations overlap (no clear winner)?

# Adding a Cost Model

---

## Definition

We translate performance metrics into monetary (or utility) costs to support design decisions.

- ▶ Basic ingredients:
  - ▶ Server cost  $c_{\text{server}}$  per server per unit time.
  - ▶ Waiting cost  $c_{\text{wait}}$  per unit waiting time per job.
  - ▶ Optional penalty for SLA violations (e.g. if  $W_q > w^*$ ).
- ▶ Given arrival rate  $\lambda$  and mean waiting time  $W_q(c)$  for configuration  $c$ :

$$C(c) = c_{\text{server}} \cdot c + c_{\text{wait}} \cdot \lambda W_q(c).$$

- ▶  $W_q(c)$  is estimated via simulation; hence  $C(c)$  is also stochastic with a CI.

# Simulation-Based Optimisation

---

- ▶ For small discrete design spaces (e.g.  $c = 1, \dots, 6$ ), we can simply enumerate options:
  - ▶ For each  $c$ , estimate  $W_q(c)$  via multiple replications.
  - ▶ Compute  $\hat{C}(c)$  and a 95% CI using the CI for  $W_q(c)$ .
- ▶ Choose the configuration with the smallest  $\hat{C}(c)$ , but report uncertainty:
  - ▶ If CIs for  $C(c)$  overlap, it may be safer to report a set of near-optimal options.
- ▶ Optional: add constraints, e.g.  $P(W_q > w^*) \leq \alpha$  (service-level constraint), and pick the cheapest feasible  $c$ .

# From Numbers to Recommendations

---

- ▶ Translate optimisation results into plain-language recommendations:
  - ▶ “With current demand, 2 agents minimise expected cost; 3 agents only improve waiting times slightly.”
  - ▶ “If demand increases by 20%, the system becomes unstable with 2 agents; we recommend adding a third.”
- ▶ Emphasise uncertainty: use CIs and scenario analysis to show robustness.
- ▶ Document assumptions: model structure, stationarity, independence, stability conditions.

# Course Pipeline Recap

---

- ▶ **Model:** choose a queueing framework and assumptions informed by data and context.
- ▶ **Estimate:** infer parameters from data (Lecture 4) and quantify input uncertainty.
- ▶ **Simulate:** implement discrete-event simulations with error control (Lecture 3).
- ▶ **Evaluate:** compare model predictions to observed performance (Lecture 5).
- ▶ **Optimise:** use costs and constraints to recommend configurations (Lecture 6).

*This is the workflow you should apply in your final projects.*

# For Your Projects

---

- ▶ Clearly state:
  - ▶ Modelling assumptions and chosen queueing framework.
  - ▶ How parameters were estimated and with what uncertainty.
  - ▶ How simulation was used to compare policies or configurations.
  - ▶ What decision or recommendation you make and how robust it is.
- ▶ Focus on a coherent modelling story rather than perfectly tuned parameters.